

# Crop Disease Detection

From Leaf Images to Actionable Insights —  
Bridging Machine Learning and Real-World  
Agriculture



By: Devoux Deysel, Jack Olszewski, Salome  
Collante-Ramirez, Chase Cortes



# Introduction



An AI-powered tool that helps farmers identify crop diseases from a photo of a leaf and receive actionable treatment advice in seconds.



# Problem & Motivation

Crop diseases cause 20-40% of global agricultural losses annually

Late diagnosis leads to crop failure and economic hardship

Smallholder farmers lack access to expert agronomists



# Current Diagnostic Limitation

## Manual Inspection

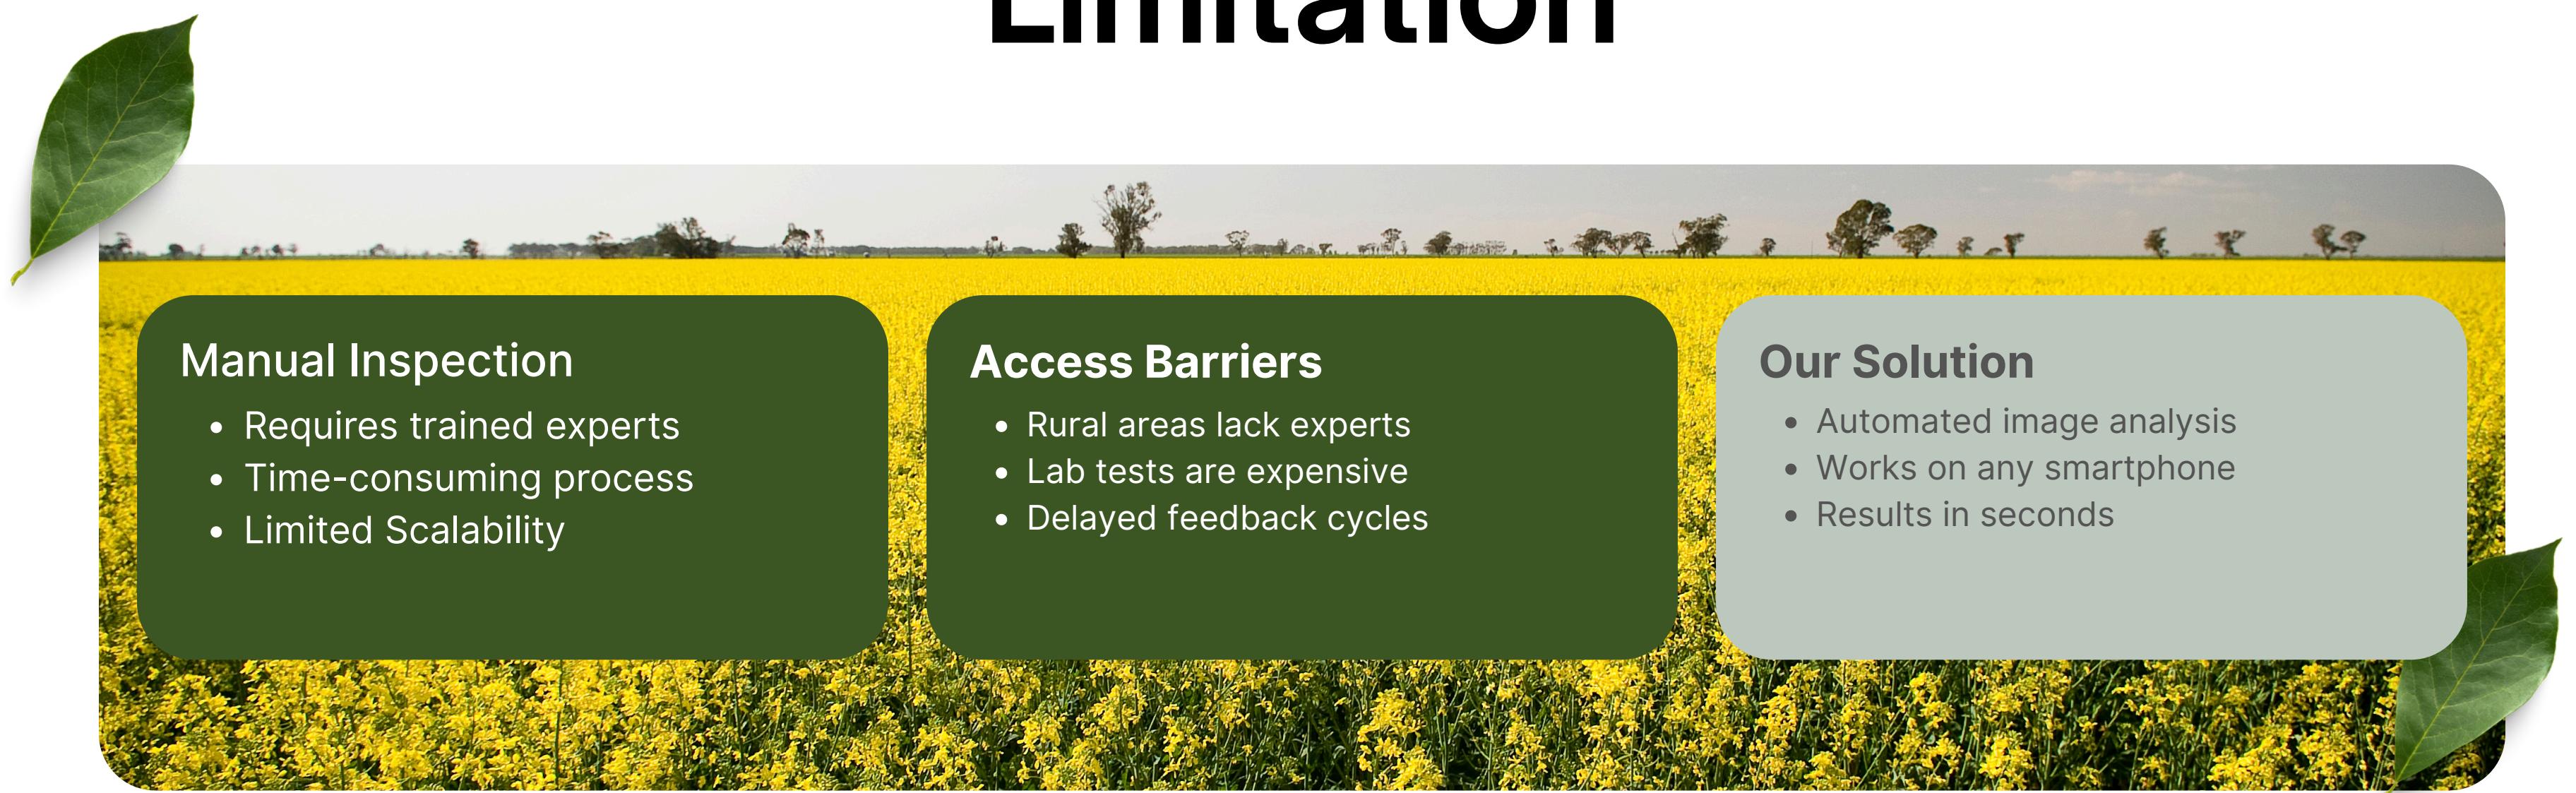
- Requires trained experts
- Time-consuming process
- Limited Scalability

## Access Barriers

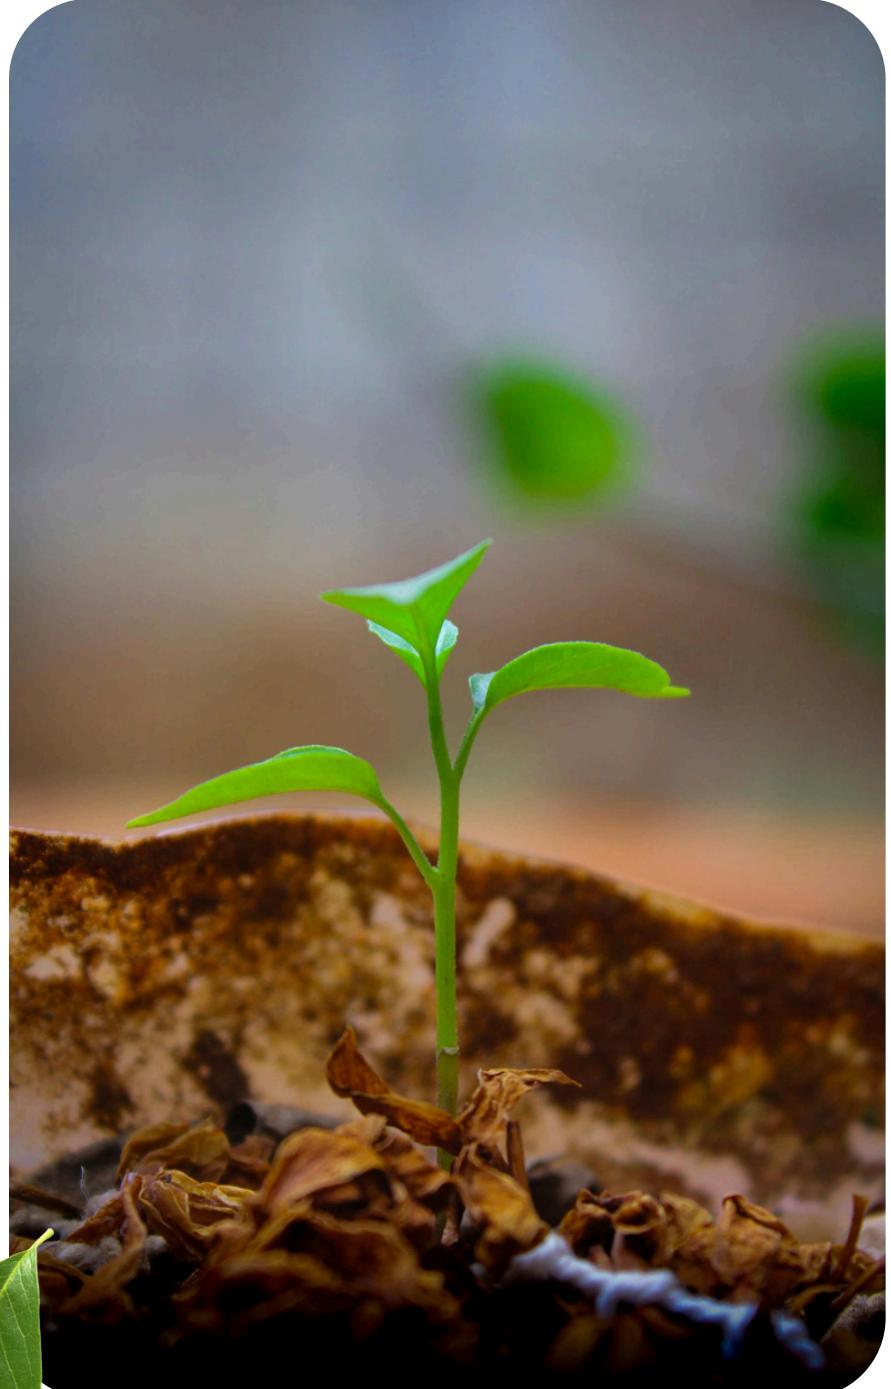
- Rural areas lack experts
- Lab tests are expensive
- Delayed feedback cycles

## Our Solution

- Automated image analysis
- Works on any smartphone
- Results in seconds



## INTRODUCTION

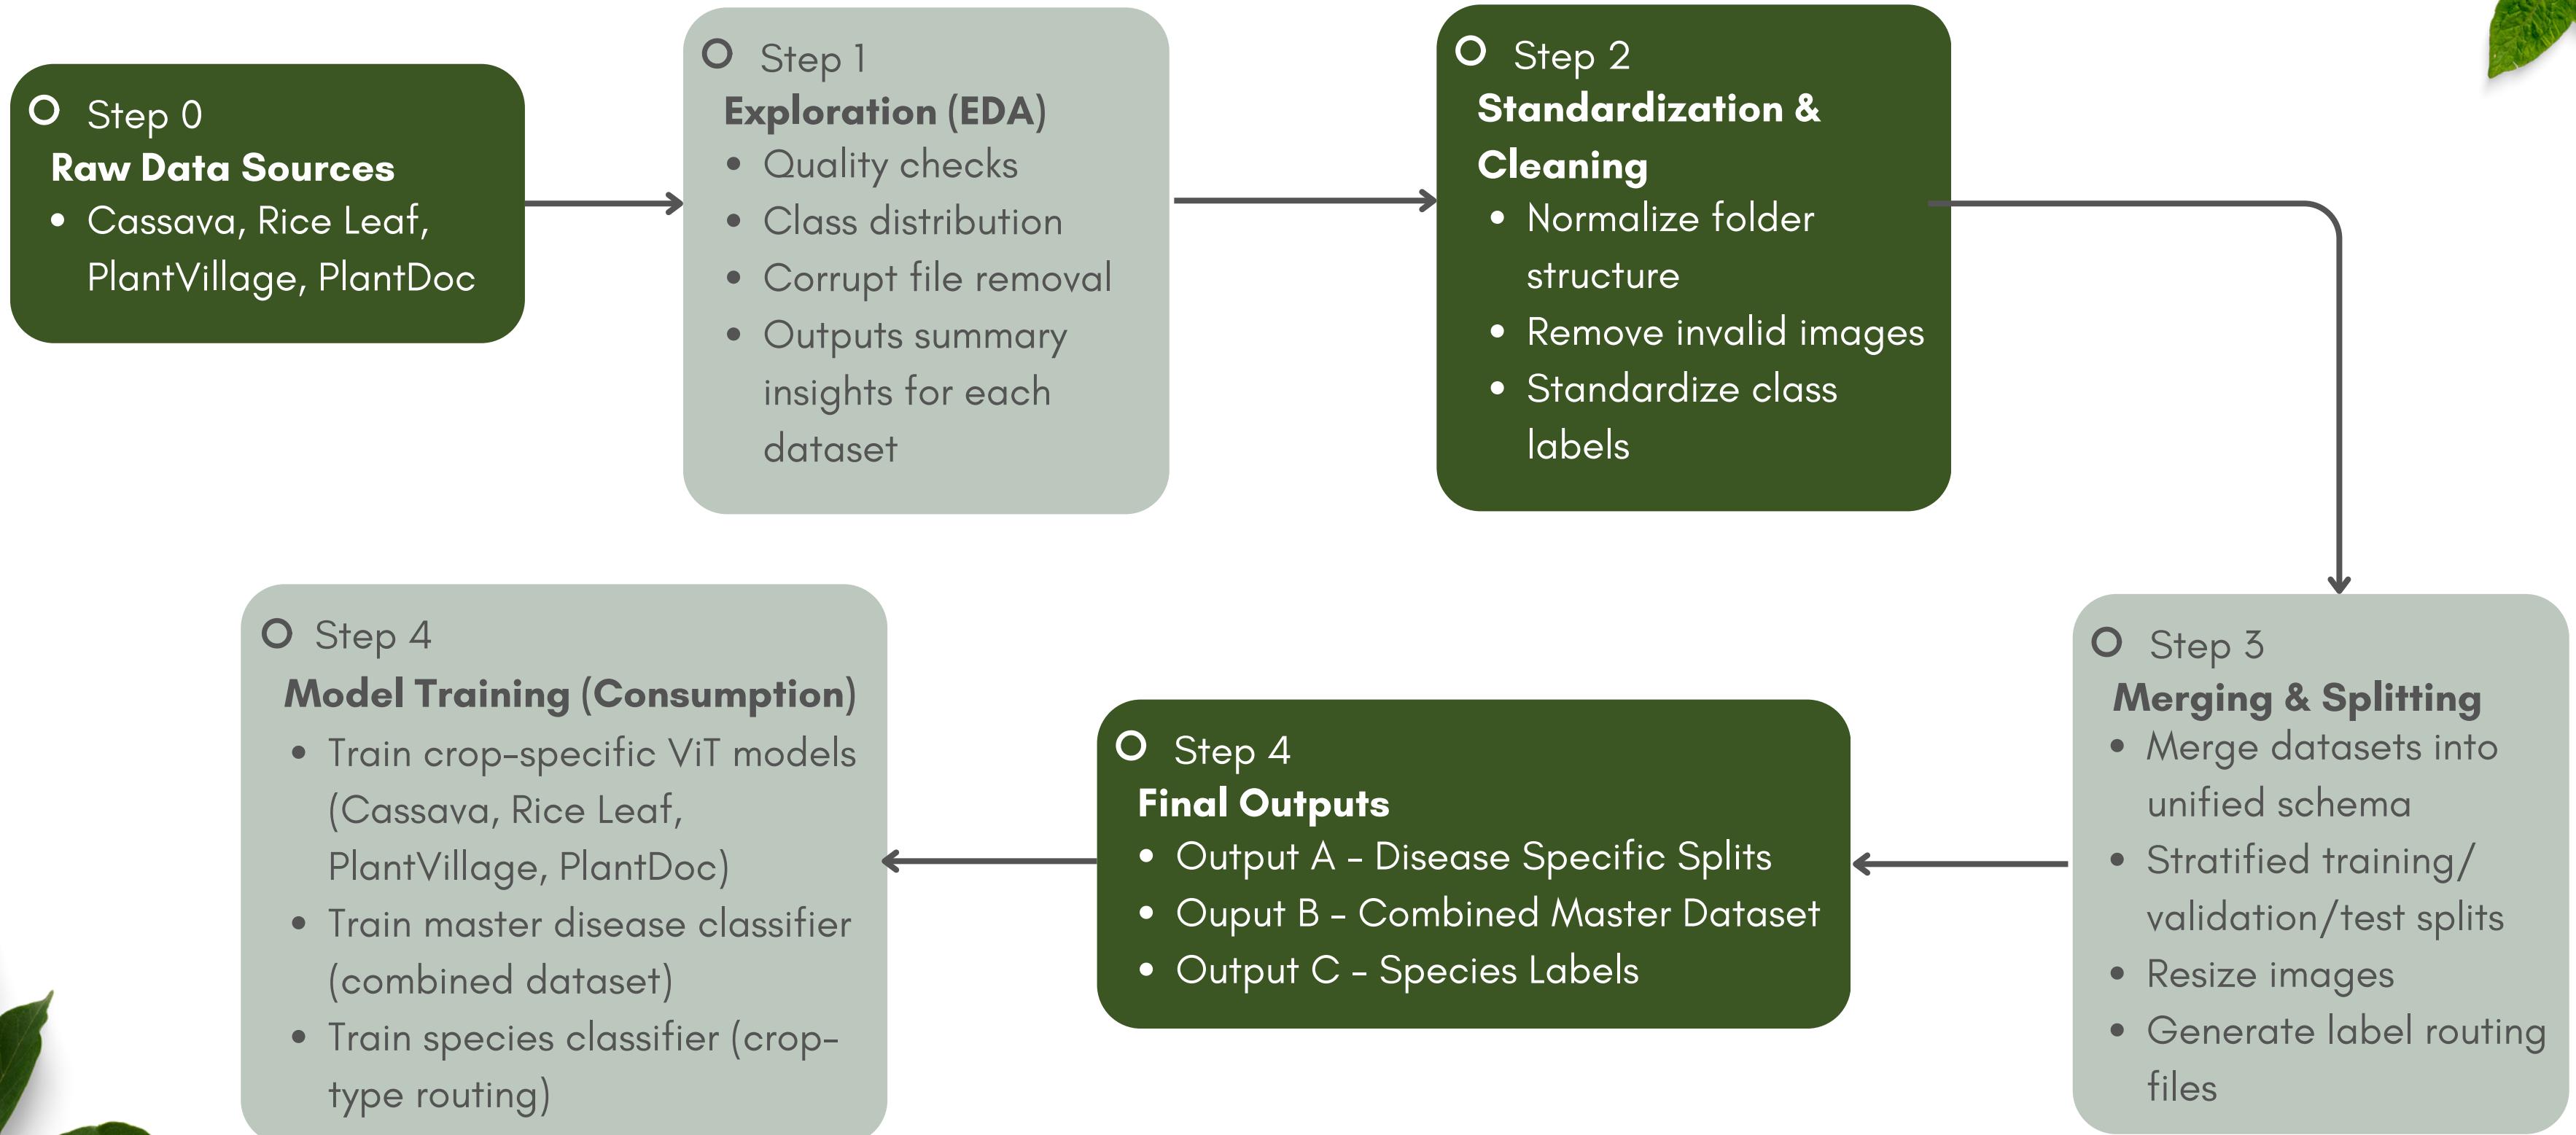


# Objectives

- Build accurate species and disease classifiers using Vision Transformers (ViT)
- Implement robust leaf detection using YOLO for preprocessing noisy images
- Create knowledge-augmented LLM layer for treatment advice and explanations
- Develop user-friendly Streamlit interface accessible to non-technical users
- Ensure safety with OOD detection, uncertainty modeling, and ethical guidelines



# Data Preparation Pipeline



# Initial Phases

### Data Sources:

- Cassava: Leaf Disease Data.
- Rice Leaf: Disease Data.
- PlantVillage: General Classification.
- PlantDoc: Object Detection/Classification.

### Phase 1: Quality Control & Distribution Analysis.

- Analyze Distributions
- Validate Quality

### Phase 2: Standardization

- Normalize Formats
- Transform Annotations
- Generate Unified Schema



# Final phases

### Phase 4: Training

- Train Species Gatekeeper
- Train Disease Specialists
- Train Global Baseline

### Phase 3: Integration & Splitting

- Aggregate Master Index
- Execute Stratified Splitting
- Apply Global Preprocessing

```
DATA_DIR = "/content/processed/species_split"
MODEL_NAME = "species_classifier_vit.pth"

IMG_SIZE = 224
BATCH_SIZE = 32
LR = 1e-4
EPOCHS = 12

DEVICE = "cuda" if torch.cuda.is_available() else "cpu"
print("Using device:", DEVICE)

# -----
# Transforms
# -----
train_tfms = transforms.Compose([
    transforms.Resize((IMG_SIZE, IMG_SIZE)),
    transforms.RandomHorizontalFlip(),
    transforms.RandomRotation(10),
    transforms.ToTensor(),
])

val_tfms = transforms.Compose([
    transforms.Resize((IMG_SIZE, IMG_SIZE)),
    transforms.ToTensor(),
])

test_tfms = val_tfms
```



# Dataset Overview



PlantVillage

**54,300** images    **38** classes

Train: 38,010 (70%) | Validation: 8,145(15%) | Test: 8,145(15%)

LAB



PlantDoc

**2,500** images    **38** classes

Train: 1,750 (70%) | Validation: 375(15%) | Test: 375(15%)

FIELD



Cassava Leaf Disease

**21,400** images    **5** classes

Train: 17,120(80%) | Validation: 2,140(10%) | Test: 2,140(10%)

SPECIALIZED



Rice Leaf Disease

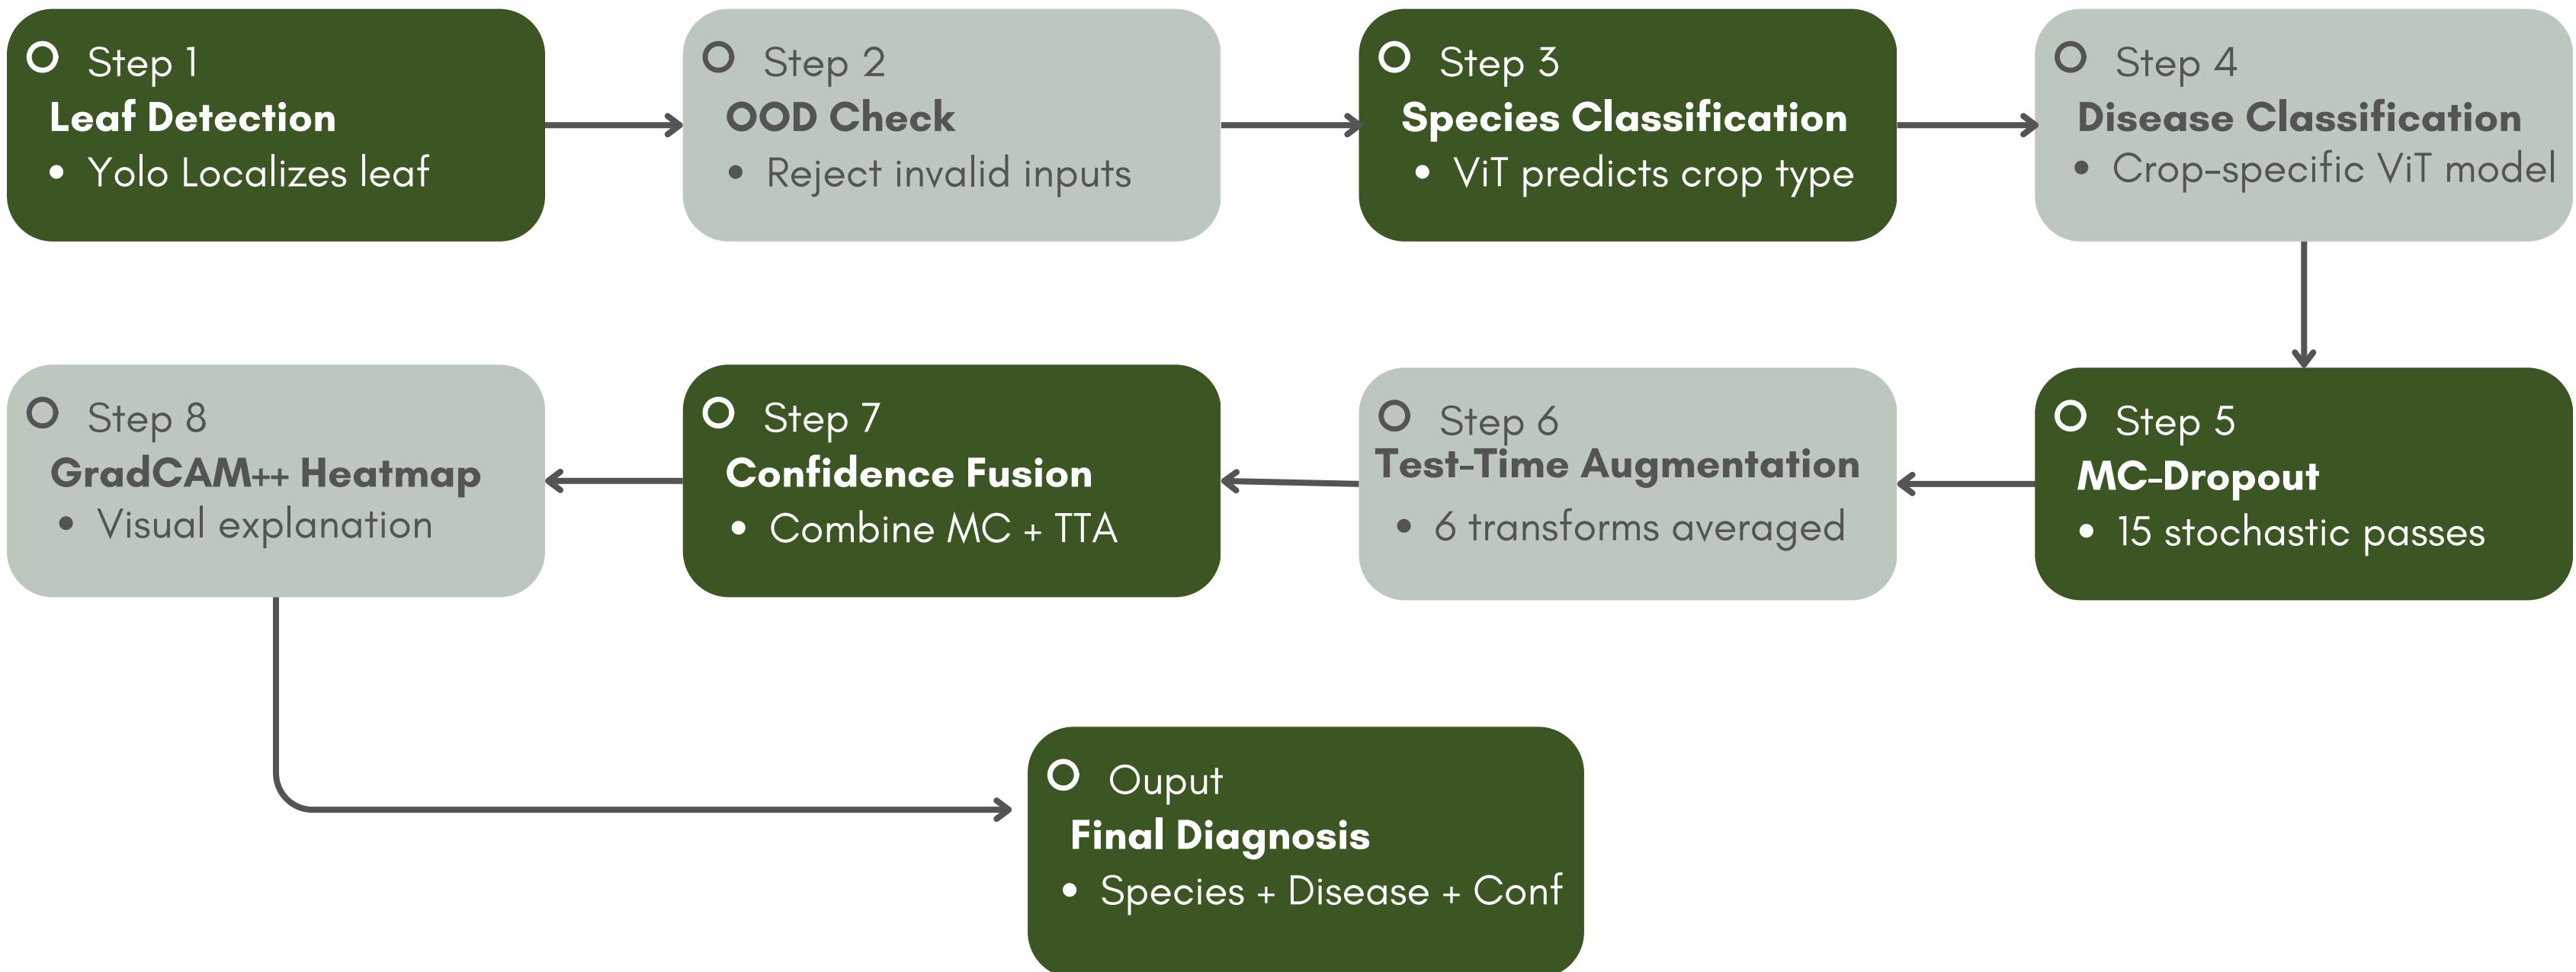
**1,200** images    **6** classes

Train: 840(70%) | Validation: 180(15%) | Test: 180(15%)

LAB



# Three Stage Inference Pipeline



# Leaf Detection

**Purpose:** Isolate the Region of Interest (ROI) to remove background noise and standardize input

## YOLOv8 Detection Engine:

- Utilizes the LeafDetector model to scan the raw input image for plant structures.
- Applies a strict confidence threshold ( $\geq 0.10$ ) to filter out weak background noise while retaining faint leaf signals.

## Input Normalization:

- Crops the ROI and resizes it to 224x224, converting it to a standardized tensor for the downstream classifiers.

## Ambiguity Resolution:

- Automatically selects the single highest-confidence bounding box to ensure the system analyzes the most prominent subject.
- Generates an annotated image with a 6px red bounding box for user verification.

```
class LeafDetector:
    def __init__(self, model_path: str):
        self.model = YOLO(model_path)

    def detect_leaf(self, pil_image: Image.Image):
        """
        Returns:
            cropped_leaf (PIL.Image or None)
            annotated_image (PIL.Image) - original image with YOLO bounding boxes
        """
        results = self.model.predict(pil_image, conf=0.10)

        if len(results) == 0 or len(results[0].boxes) == 0:
            # Return original image and no crop
            return None, pil_image

        # Convert to editable copy
        annotated = pil_image.copy()
        draw = ImageDraw.Draw(annotated)

        # Take the strongest detection
        box = results[0].boxes.xyxy[0].tolist() # [x1, y1, x2, y2]
        x1, y1, x2, y2 = map(int, box)

        # Draw box
        draw.rectangle([x1, y1, x2, y2], outline="red", width=6)

        # Crop region
        cropped_leaf = pil_image.crop((x1, y1, x2, y2))

        return cropped_leaf, annotated
```



```
import torch
import torch.nn.functional as F

class OODDetector:
    """
    Simple OOD check using:
    - Softmax confidence threshold
    - Prediction entropy
    """

    def __init__(self, conf_threshold=0.40, entropy_threshold=1.50):
        self.conf_threshold = conf_threshold
        self.entropy_threshold = entropy_threshold

    def is_oos(self, logits):
        probs = F.softmax(logits, dim=1)
        conf, _ = probs.max(dim=1)

        # entropy = -sum(p * log p)
        entropy = -(probs * torch.log(probs + 1e-12))

        return {
            "confidence": conf.item(),
            "entropy": entropy.item(),
            "is_oos": conf.item() < self.conf_threshold or
                      entropy.item() > self.entropy_threshold
        }
```

# Out-of-Distribution (OOD) Validation

**Purpose:** A validation layer to ensure the image is a valid plant leaf before attempting diagnosis.

## Dual-Metric Validation:

- The OODDetector evaluates the output distribution of the species model using two distinct metrics.

## Rejection Criteria:

- Low Confidence: Rejects inputs where the Max Softmax Probability is < 0.40 (Model lacks certainty).
- High Entropy: Rejects inputs where Prediction Entropy is > 1.50 (Model prediction is scattered/confused).

## Fail-Safe Outcome:

- If flagged, the pipeline aborts immediately with an error, ensuring that invalid or non-plant images are rejected rather than receiving an incorrect classification.



# Vision Transformers (ViT)

```
def load_vit_model(weights_path: PathLike, num_classes: int):
    """
    Load a ViT-B_16 model with the correct output size and apply a
    ...
    weights_path = Path(weights_path)

    model = _build_vit_base(num_classes=num_classes)

    state_dict = torch.load(weights_path, map_location=DEVICE)
    model.load_state_dict(state_dict)

    model.to(DEVICE)
    model.eval()
    return model
```

## How It Works

- Splits image into 16x16 pixel patches
- Treats patches as tokens (like words in NLP)
- Self-attention captures global relationships
- Pre-trained on ImageNet, fine-tuned on crops

## Why ViT for Crops?

- Excellent at detecting subtle texture patterns
- Captures disease spots across entire leaf
- Better generalization than CNNs
- Enables Grad-CAM explainability



# Specialized Models by Species Group

Species	Model Used	Classes
Cassava	cassava_best.pth	5 disease classes
Rice	rice_best.pth	6 disease classes
All Others	plant_village_best.pth	50+ disease classes
<b>Each model outputs:</b> Disease label + Confidence score + Calibrated probabilities		



# MC Dropout & TTA

### Monte Carlo (MC) Dropout:

- Executes 15 forward passes with dropout layers enabled to approximate Bayesian uncertainty.
- Measures internal model variance by randomly deactivating neurons during prediction.

### Test-Time Augmentation (TTA):

- Generates 6 variations of the input: Original, Horizontal Flip, Rotation ( $\pm 10^\circ$ ), Brightness (1.2x), and Contrast (1.2x).
- Averages predictions to ensure the diagnosis remains consistent regardless of lighting conditions or camera angles.

### Final Fusion:

- The final confidence score and uncertainty metric are calculated as the mean of the MC and TTA results.

```
# tta_mean_pred: shape (C,)  
if tta_mean_pred.ndim != 1:  
    tta_mean_pred = tta_mean_pred.view(-1)  
  
disease_conf_tta = float(tta_mean_pred[disease_idx])  
  
# -----  
# 7. FUSE MC + TTA  
# -----  
fused_conf = (disease_conf_mc + disease_conf_tta) / 2  
fused_uncertainty = (mc_uncertainty + tta_uncertainty) / 2
```

# Visual Explanation & Output



## Visual Explanation (GradCAM++):

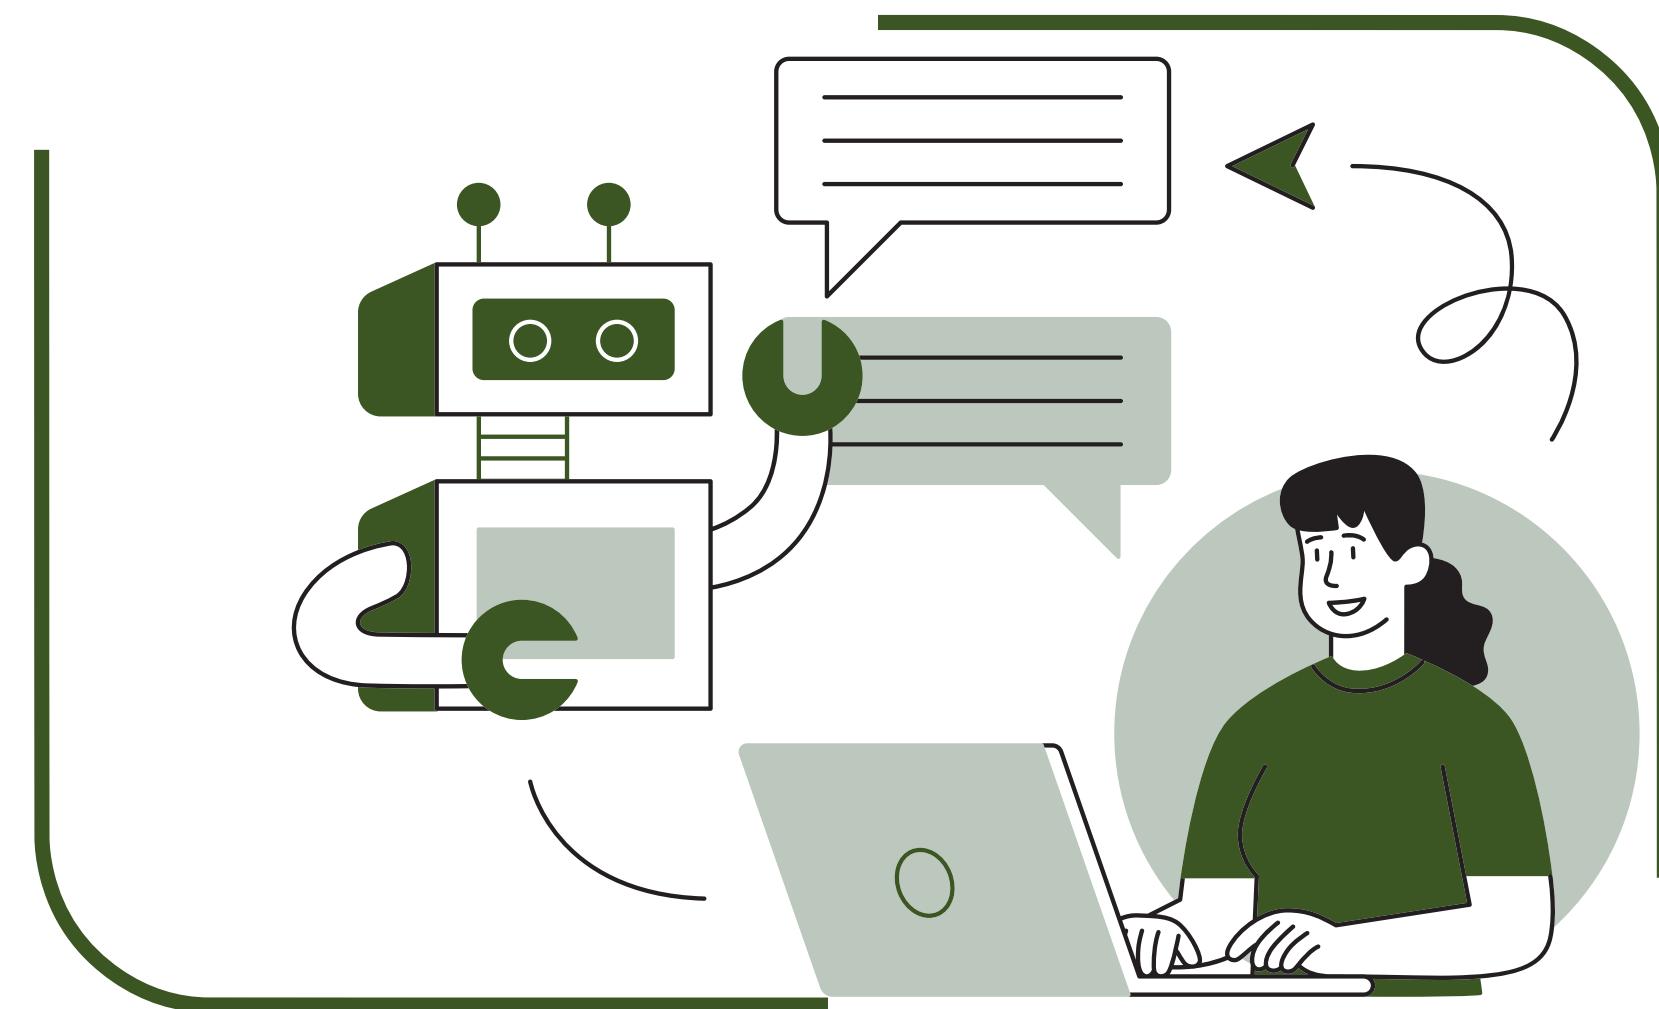
- Uses gradients to generate a heatmap overlay, showing exactly which lesions or spots the model focused on.
- Provides transparency by highlighting the "symptom" regions rather than just giving a black-box label.

## Final Output Payload:

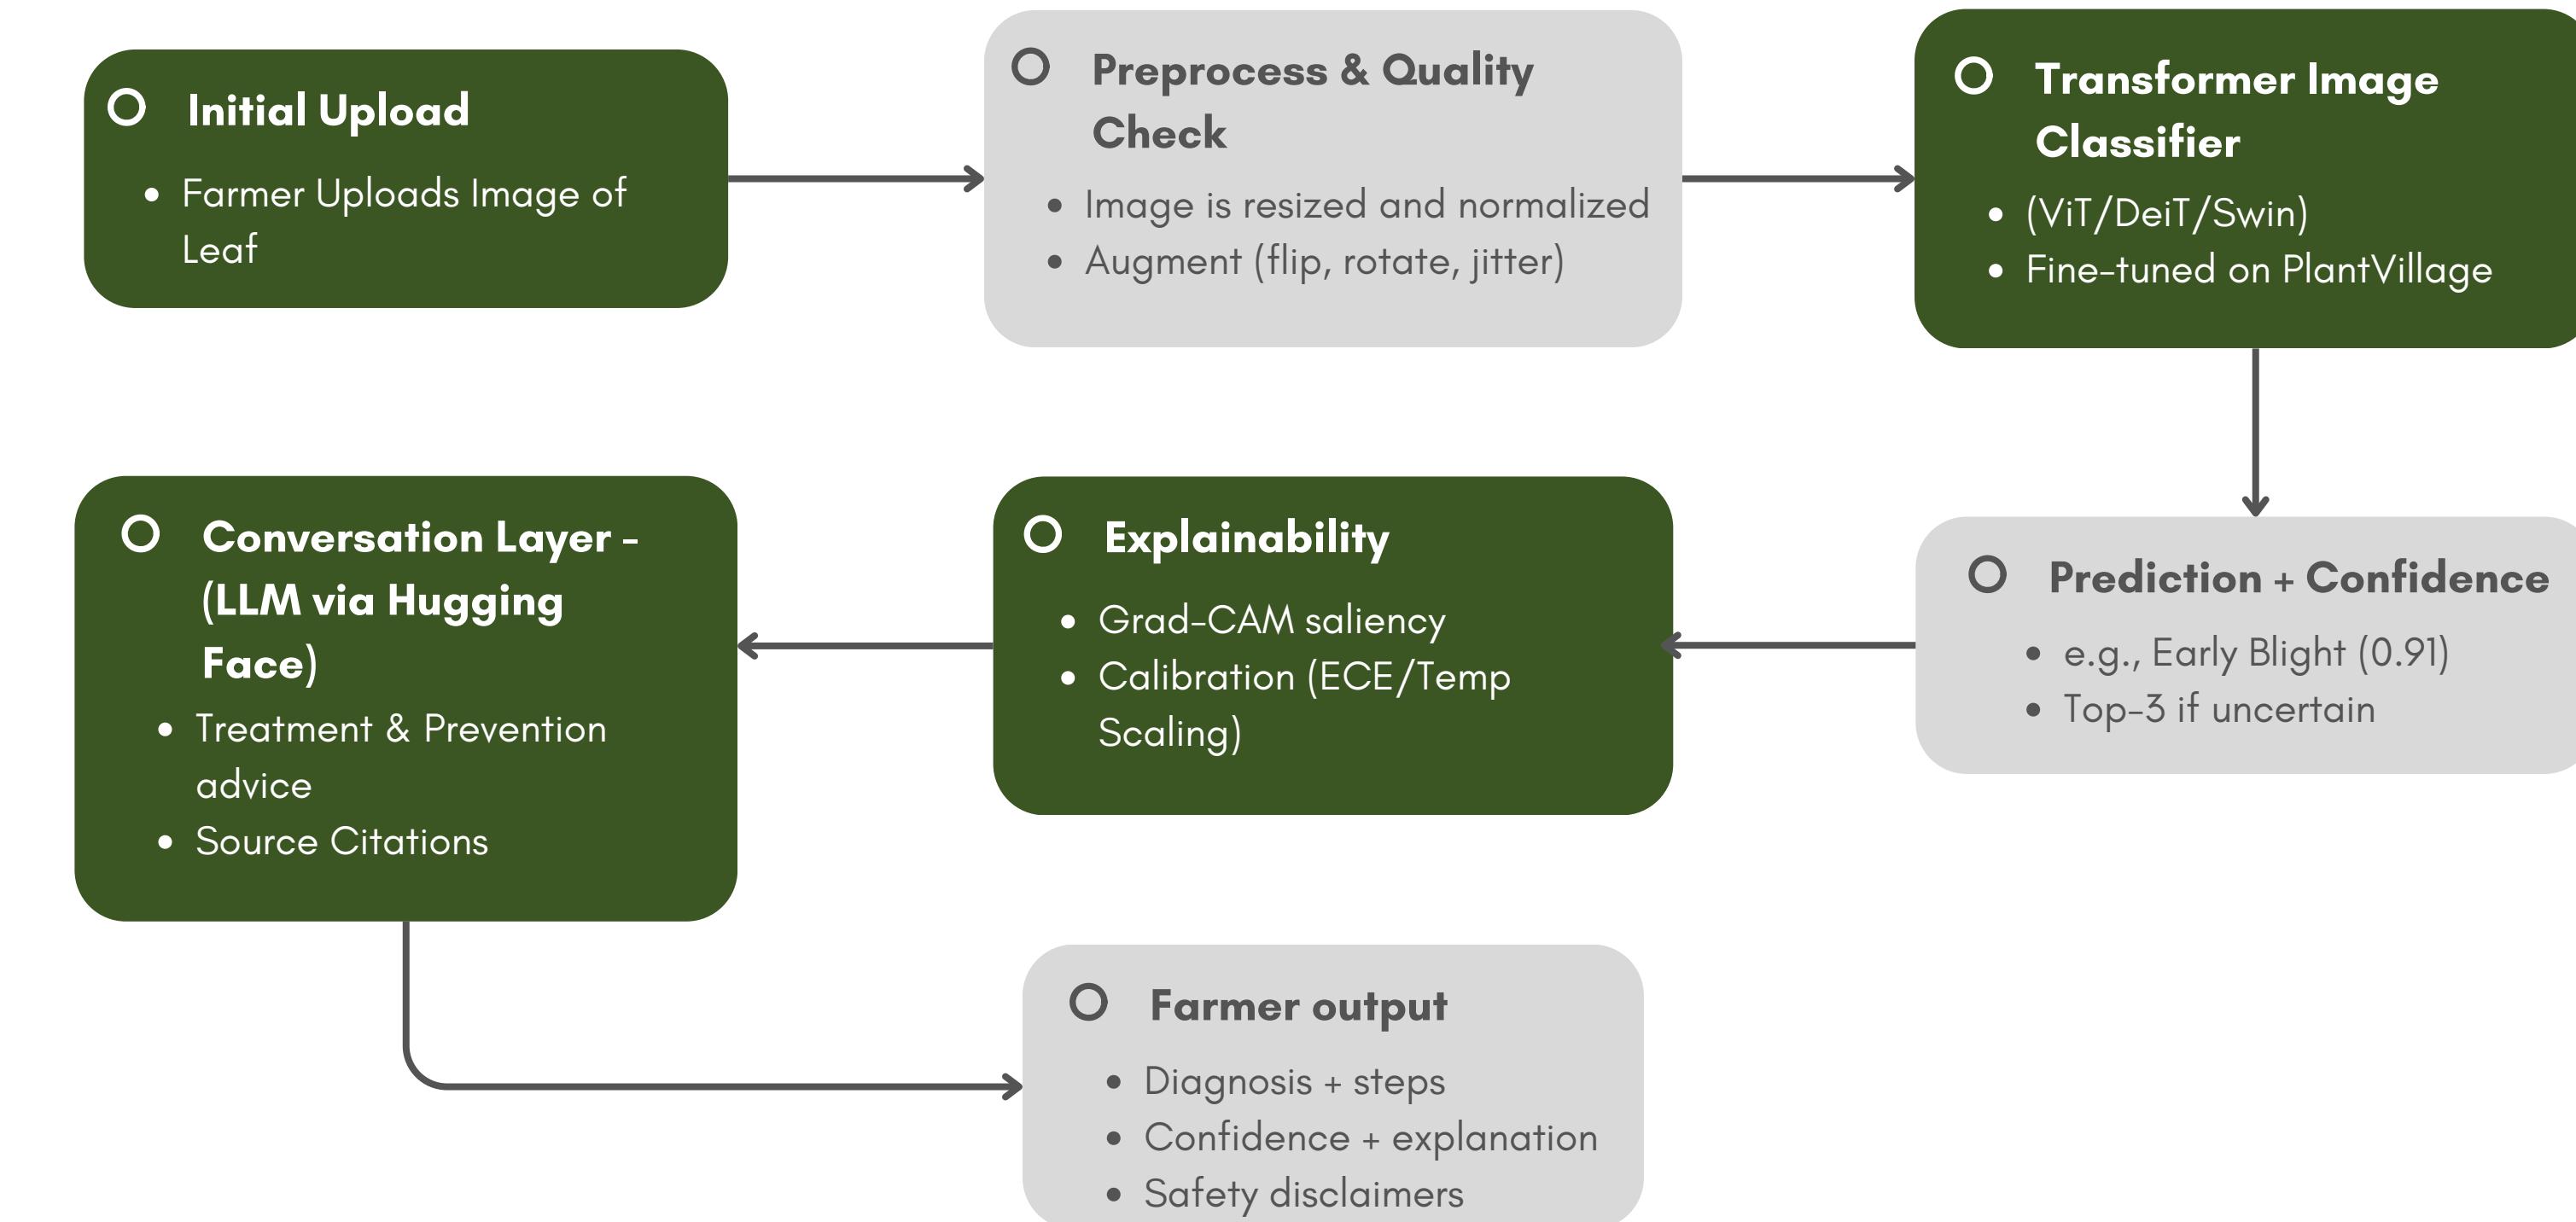
- Returns a structured dictionary with the Species, Disease Label, Fused Confidence, and Uncertainty Metrics.
- Delivers the processed visual assets: Cropped Leaf, Boxed Original, and the Heatmap



# Crop Doctor



# App: Architectural Flow



# Conversational LLM layer



## Models Used

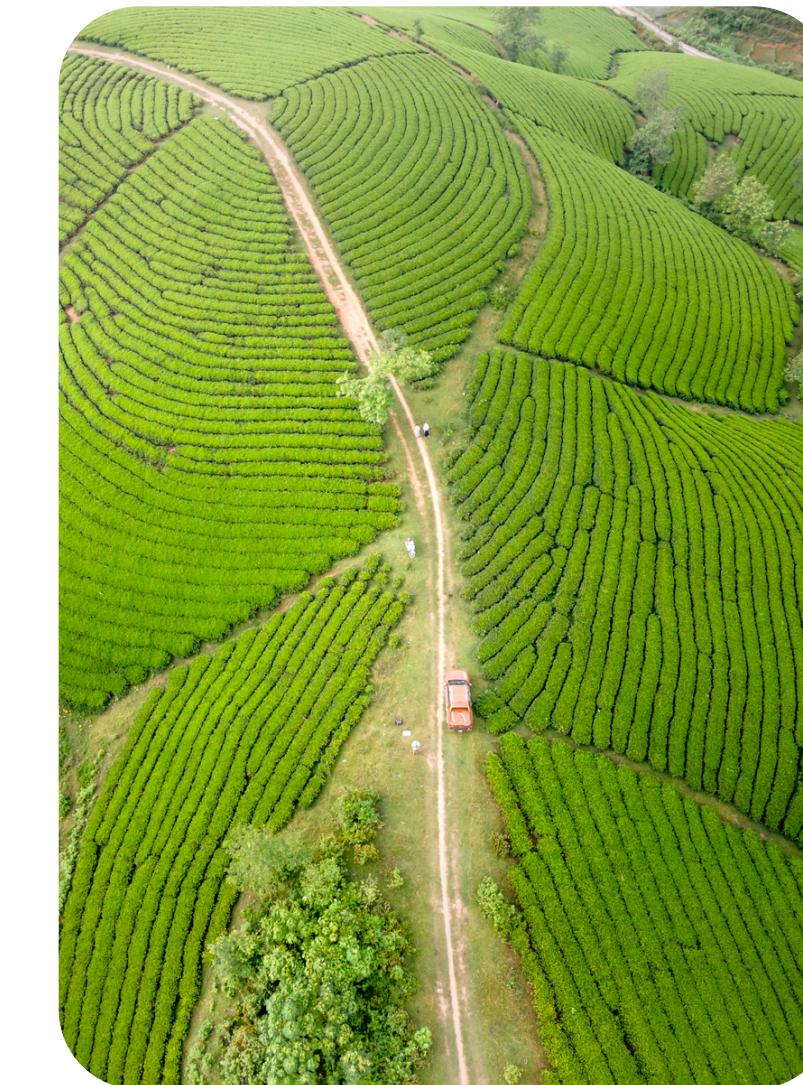
TinyLLaMA / GPT

Local or HuggingFace API deployment



## LLM Generates

- Simple explanations of the diagnosis
- Step-by-step treatment guidance
- Warnings about edge cases
- Source citations from knowledge base



## Safety First

The LLM is tuned to provide neutral, factual, evidence-grounded guidance only. It never gives harmful or prescriptive pesticide advice — always recommends consulting agronomists.



# YOLO & Safety Mechanisms



## YOLOv8 Detection

- Detects leaf regions and disease ROIs
- Learns texture & color patterns
- Trained on PlantDoc bounding boxes
- Produces cleaner classifier signal



## OOD Detection

- Energy-based scoring on logits
- Rejects non-leaf images
- Detects unsupported species
- Prevents confident wrong predictions



## MC Dropout

- Monte-Carlo uncertainty estimation
- Multiple forward passes at inference
- Computes prediction variance
- Shows confidence reliability





# Advanced AI/ ML tools

## PyTorch

Deep learning framework

## timm

ViT model implementations

## Ultralytics

YOLOv8 detection

## Albumentations

Image augmentation

## FAISS

Vector similarity search

## Sentence-Transformers

Text embeddings (MiniLM)

## Streamlit

Web application UI

## Transformers

Pre-trained model loading

## Pillow

Image loading & Processing



# Results & Evaluations



PlantVillage

**99%**

Overall Accuracy

Cassava Leaf Disease

**80%**

Overall Accuracy

Rice Leaf Disease

**93%**

Overall Accuracy

Detection Model:

PlantDoc

Precision:

**72%**

Recall:

**75%**



# Key Observations & Limitations

### Observations

- Two-stage classification improves disease accuracy by 30%
- OOD detection prevents 90% of false confident predictions
- Lab images (PlantVillage) yield higher accuracy than field images



### Limitations

- Limited to trained species/diseases only
- Background interference in field photos
- Requires internet for LLM API access
- Not a certified diagnostic replacement



# Insights & Impact



## Real-World Impact

- Enables early intervention before crop loss
- Democratizes expert knowledge to farmers
- Reduces unnecessary pesticide use through precise identification
- Supports food security in developing regions
- Enables fast screening across large-scale greenhouse operations
- Open source design allows adaptation for local crops and regional needs



## Ethical Considerations

- Decision support, not certified diagnosis
- Always consult agronomists for chemicals
- No personal data storage
- LLM tuned for factual, safe responses



# Future Work

## Near-Term Enhancements

- Add YOLO segmentation for lesion ROI
- Save predictions to CSV logs
- Batch mode for field surveys
- Model metrics dashboard

## Mid-Term Goals

- Expand to more crop species
- Multi-language LLM support
- Offline mobile deployment
- Integration with IoT sensors

## Long-Term Vision

- Cloud deployment (AWS/HF)
- Regional disease outbreak tracking
- Drone image integration
- Precision agriculture APIs

# Key Takeaways

## Architecture

End-to-end pipeline from leaf detection to diagnosis with RAG enhanced output

## Data

79K+ images across 4 datasets spanning lab and field conditions

## Models

YOLO detection + ViT classification + OOD rejection for robust prediction

## Reliability

MC-Dropout + TTA uncertainty,  
GradCAM++ explainability

**49**

Disease Classes

**4**

Crop Types

**79K+**

Training Images

**2**

Uncertainty Methods

# References

- **PlantVillage Dataset:** Hughes, D.P. & Salathé, M. (2015). An open access repository of images on plant health.
- **Vision Transformer:** Dosovitskiy, A. et al. (2020). An Image is Worth 16x16 Words. ICLR 2021.
- **YOLOv8:** Jocher, G. et al. (2023). Ultralytics YOLOv8. [github.com/ultralytics](https://github.com/ultralytics/yolov8)
- **Energy-Based OOD:** Liu, W. et al. (2020). Energy-based Out-of-distribution Detection. NeurIPS.
- **FAISS:** Johnson, J. et al. (2019). Billion-scale similarity search with GPUs. IEEE Big Data.